

Distributed SPMV algorithm design and performance analysis

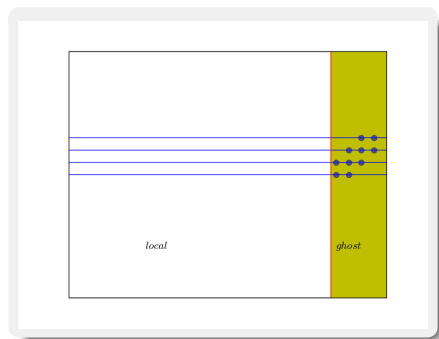
Zhengjiang Li

Department of Mechanical and Aerospace Engineering
University at Buffalo, State University of New York

June 15, 2015

Data Structure

- Distributed Vector
- Distributed Sparse Matrix




Distributed SPMV

Algorithm 1 SPMV

```


1: for  $i = 0; i < \text{LocalNumberOfRows}; i++$  do
2:    $\text{cur\_mat\_val} = \text{matrixValues}[i]$ 
3:    $\text{ur\_local\_ind} = \text{matrixInd}[i]$ 
4:    $\text{cur\_nnz} = \text{nnzInRow}[i]$ 
5:   for  $j = 0; j < \text{cur\_nnz}; j++$  do
6:      $\text{sum} += \text{cur\_mat\_val}[j] * \text{vec.val}[\text{cur\_nnz}[j]]$ 
7:   end for
8: end for
    
```

spmv



Communication Buffer

GlobalRow	nnzCollnCurRow	Interface Communicate
101	(96, 97, 101, 102, 106, 107)	(j, 106, 107)
102	(96, 97, 98, 101, 102, 103, 106, 107, 108)	(j, 106, 107, 108)
103	(97, 98, 99, 102, 103, 104, 107, 108, 109)	(j, 107, 108, 109)
...



CG algorithm

Algorithm 2 CG

```

SparseMatrix A
Vector x0, Ax0, b, r, d
spmv(A, x0)
ScalarProduct(b, Ax0, r) % initial residual
DotProduct(r, r, normr); % residual norm
norm0 = norm
for i = 1; i < max_iter && normr/norm0 > eps; i++ do
    beta = norm/norm0;
    ScalarProduct(z, beta, d, d) %update direction vector
    spmv(A, d, Ad)
    DotProduct(d, Ad, dAd);
    alpha = norm/dAd;
    ScalarProduct(x, alpha, d, x); %update solution vector
    ScalarProduct(r, -alpha, Ad, r); %update residual vector
    DotProduct(r, r, normr)
end for

```

Time Complexity

topology	localX	localY	buffer
1D	NX	NY/p	2NX
2D	NX/px	NY/py	2(NX/px + NY/py)
...	

CPU calculation time :

$$t_1 = (2m + 2N)/c_1 = 2(nnzInRow + 1)N/c_1 \quad (1)$$

Communicate time:

$$t_2 = 2NX/c_2 \% 1D \text{ cpu} \quad (2)$$

$$t_2 = 2(NX/px + NY/py)/c_2 \% 2D \text{ cpu} \quad (3)$$

Time Complexity II

$$t = t_1 + t_2 = \begin{cases} 2\delta \frac{NX \cdot NY}{p \cdot c_1} + 2 \frac{NX}{c_2} \% 1D \text{ cpu} \\ 2\delta \frac{NX \cdot NY}{p \cdot c_1} + 2 \frac{NX \cdot py + NY \cdot px}{px \cdot py \cdot c_2} \% 2D \text{ cpu} \end{cases}$$

The time complexity is proportional to the size of problem, nonzeros in rows, as well as hardware performance. The speed-up will reach an upper limit when:

- ① For 1D CPU topology, $p \geq \gamma NX$
- ② For 2D CPU topology, $2\sqrt{p} \geq \gamma NX$

Also consider accuracy(iteration numbers)

Results





Figure : CPU performance

sparse matrix src

```

1 struct SparseMatrix
2 {
3     Geometry*   geom;
4
5     /* local variable */
6     local_int_t localNumberOfRows;
7     local_int_t localNumberOfColumns;
8     local_int_t localNumberOfNonzeros;
9     global_int_t** mtxIndG;
10    global_int_t** mtxIndL;
11    double** matrixValues;
12
13    global_int_t* nonzerosInRow;
14
15    std::map<global_int_t, local_int_t> globalToLocalMap;
16    std::vector<global_int_t> localToGlobalMap;
17

```

Thank you for your attention

1

¹ Jonathan Richard Shewchuk "An introduction to the Conjugate Gradietn Metho